

Assignment Project Exam Help

ST227: Applications of R in Life Insurance
Optimisation in MLE & Applications in Markovian Models

<https://powcoder.com>

Dr. Viet Dang

Add WeChat ^{21/02/2021} powcoder

- ▶ Often in mathematical sciences, one is concerned with the following question:

given an input domain \mathcal{O} and a criterion function $f : \mathcal{O} \rightarrow \mathbb{R}$,

- ▶ Is there a minimum point $x \in \mathcal{O}$?
- ▶ Is this explicitly solvable? If it's not, can we efficiently find it numerically?
- ▶ The nature of optimal point: min? max? inflection? local? global?
- ▶ If \mathcal{O} is a subset of \mathbb{R}^n , then this procedure often involves:
 - ▶ Solving for the candidate solution, which satisfies:
$$\nabla f = 0 \quad i = 1, \dots, n.$$
 - ▶ Verify the nature of this point by consider the positive or negative-definiteness of the Hessian matrix:

$$D^2 f \triangleq \begin{pmatrix} \partial_{11}^2 f & \partial_{12}^2 f & \dots & \partial_{1n}^2 f \\ \partial_{21}^2 f & \partial_{22}^2 f & \dots & \partial_{2n}^2 f \\ \vdots & \vdots & \ddots & \vdots \\ \partial_{n1}^2 f & \partial_{n2}^2 f & \dots & \partial_{nn}^2 f \end{pmatrix}$$

Assignment Project Exam Help

- Let us consider a toy example where $f: x \mapsto x^2$. Then

$$f'(x) = 2x, \quad x = 0$$

$$f''(x) = 2, \quad f''(x)|_{x=0} = 2 > 0.$$

The candidate optimiser is $x = 0$. The second derivative test implies that it is a local minimum.

- ▶ We **cannot** conclude that it is a global minimum without further considerations.
- ▶ If the second derivative test is 0, we cannot make a conclusion.
- ▶ Example: $g(x) = x^4$, $h(x) = -x^4$ and $k(x) = x^3$ all have $x = 0$ as the unique candidate - yet the nature of this point is different.

Assignment Project Exam Help

- ▶ In practice – it is extremely difficult to solve $f' = 0$ analytically.
- ▶ Example: $\mathcal{O} = (0, \infty)$ and $f : x \rightarrow \frac{1}{2}x^2 \ln(x) - \frac{1}{4}x^2 - 2x$. It has a (exercise!) unique global minimiser. The first order condition is:

$$f'(x) = x \ln(x) - 2 = 0.$$

<https://powcoder.com>

Can you solve it “by hand”?

- ▶ Potential numerical approaches:
 - ▶ Interval Bisection
 - ▶ Newton-Raphson algorithm
- ▶ You can write your own algorithm or ...

Add WeChat powcoder

Assignment Project Exam Help

- ▶ `optim` is a general-purpose optimisation routine in R.
- ▶ General syntax: `optim(par,fn)`, where:
 - ▶ `par`: initial values of parameters;
 - ▶ `fn`: a function to be minimised.
- ▶ Additional fine-tuning is available. We will introduce them as we need them. Type `?optim` for its documentation. Some additional arguments of interest are:
 - ▶ `method`: Method of optimisation. Default value is the “Nelder-Mead” method.
 - ▶ `lower` and `upper`: bounds for the candidate solution.
- ▶ Available methods include: “Nelder-Mead”, “BFGS”, “CG”, “L-BFGS-B”, “SANN”, and “Brent”. Their algorithms are beyond the scope of ST227.

Add WeChat powcoder

- ▶ Consider the toy example: $f: x \rightarrow x^2$. Use 1 as the initial value.

```
sqr = function(x){x^2}
optim(par=1,sqr)
```

- ▶ The output is a list. We are interested in the following items: \$par = -8.881784e-16 is the minimiser, and \$value = 7.888609e-31 is the the function value at this minimiser.
- ▶ Be careful of the warning, though: "Warning in optim(par = 1, sqr): one-dimensional optimization by Nelder-Mead is unreliable: use 'Brent' or optimize() directly"
- ▶ The default method is designed for multi-dimensional optimisation. Let us use the Brent method
- ▶ The Brent method requires lower and upper arguments.

```
optim(par=1,sqr,method="Brent",lower=-10,upper=10)
```

Assignment Project Exam Help

What if a function does not have a maximum/minimum?

```
id <- function(x){x}  
optim(par=0,id)  
optim(par=0,id,method="Brent",lower=-10,upper=10)
```

Does this result make sense?

- ▶ Use numeric optimisation with reservation.
- ▶ Exercise: Use optim with method = "Brent" to minimise on $(0, \infty)$ the function:

Add WeChat powcoder

$$f: \mathbb{R} \rightarrow \mathbb{R}, x \mapsto x^2 \ln(x) - \frac{1}{4}x^2 - 2x$$

- ▶ Consider the function: $f : (x, y) \rightarrow x^2 + y^2$, which has a global minimum at the origin.

The underlying syntax of `optim` requires f to have a **single vectorial argument**, instead of two real-valued arguments. That is, instead of:

```
f <- function(x,y){  
  x^2 + y^2  
}
```

we will need:

```
f <- function(x){  
  x[1]^2 + x[2]^2  
}
```

- ▶ For a multi-dimensional optimisation problem, we can use the default Nelder-Mead method:

```
optim(par = c(1,1), f)
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

How does one choose initial values for an optimisation algorithm? There are no definitive answers. Some approaches are:

- ▶ Numerical Experimentation,
- ▶ Qualitative insight of the problem, e.g. constrained optimisation,
- ▶ Inspiration from similar, but simpler problems,
- ▶ Guesses.
- ▶ Minimisation algorithms actually search for a turning point.
 - ▶ No guarantee on the uniqueness/nature of the turning point.
 - ▶ Might get trapped in local optima.
- ▶ Definitely NOT a replacement for mathematical insights.
- ▶ Widespread use in Statistics: Maximum Likelihood Estimation, Least Square Estimation (e.g. Linear Regression and GLM), Stochastic Gradient Descent (Deep Learning, Neural Networks).

<https://powcoder.com>

Add WeChat powcoder

Maximum-Likelihood Estimation

- ▶ Suppose that we observe a sample (X_1, \dots, X_n) , which depend on a common parameter $\theta \in \mathbb{R}^d$, where $X_i \sim f_i(x, \theta)$.
- ▶ Questions: 1. How to estimate θ based on (X_1, \dots, X_n) ? and 2. How sure are we of this estimate?
- ▶ There are various approaches:
 - ▶ Method of Moments;
 - ▶ Least Square Estimation;
 - ▶ Maximum Likelihood Estimation
- ▶ The likelihood $L(\theta; X)$ and log-likelihood $l(\theta; X)$ are defined as follows:

$$L(\theta; X) = \prod_{i=1}^n f_i(x_i, \theta) \quad (1a)$$

$$l(\theta; X) = \sum_{i=1}^n \log(f_i(x_i, \theta)) \quad (1b)$$

- ▶ The maximum likelihood estimator is defined as the solution of:

$$\max_{\theta} l(\theta; X),$$

which is an optimisation problem.

- ▶ For most distributions seen in ST102, this problem is explicitly solvable. In general, though, we will have to solve it numerically.

Assignment Project Exam Help

- ▶ A random variable X follows a $\mathcal{N}(\mu, \sigma^2)$ if it has density:

$$f(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}, \quad x \in \mathbb{R}.$$

- ▶ Given an i.i.d. normal sample $X = (X_1, \dots, X_n)$, let us estimate μ and σ^2 using numerical MLE method.
- ▶ Suppose $n = 100$. We will simulated some dummy data for this task:

```
X <- rnorm(n=100, mean=1.5, sd=5)
```

Add WeChat powcoder

Assignment Project Exam Help

► Maximising log-likelihood means minimising its negative. That is, we are minimising the function:

$$\sum_{i=1}^n -\log \left(\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}} \right)$$

<https://powcoder.com>

► There are at least three ways I can think of to define this sum, with varying level of difficulty:

- 1: using for loop and elementary calculations only.
- 2: utilising R's vectorisation.
- 3: utilising R's built-in statistical functions.

Add WeChat powcoder

- ▶ The first method is a *low-level method*. If you have studied a “traditional” programming course before, this is probably how you define a sum:

```
nLL <- function(param){  
  mu <- param[1]  
  sigma2 <- param[2]  
  out <- 0  
  for(i in 1:length(X)){  
    out <- out - log((2*pi*sigma2)^(1/2)*exp(-(X[i]-mu)^2/(2*sigma2)))  
  }  
  return(out)  
}
```

- ▶ This displays all the fine details in code.
- ▶ Very reminiscent of low-level languages e.g. C.
- ▶ This is **not** recommended. Too much bookkeeping obfuscates the essential ideas.

- ▶ The second method employs R's vectorised calculations. Those of you with some Matlab experience will be familiar with this.

```
nLL <- function(param){  
  mu <- param[1]  
  sigma2 <- param[2]  
  - sum(  
    log((2*pi*sigma2)^{-1/2}*exp(-(X-mu)^2/(2*sigma2)))  
  )  
}
```

<https://powcoder.com>

- ▶ Why does this work? Due to vectorisation, the code `X-mu` actually returns:

$$(X_1 - \mu, X_2 - \mu, \dots, X_n - \mu).$$

- ▶ The code `(X-mu)^2` returns: $((X_1 - \mu)^2, \dots, (X_n - \mu)^2)$.
- ▶ Tracing along all the calculations, we see that the code `log((2*pi*sigma2)^{-1/2}*exp(-(X-mu)^2/(2*sigma2)))` returns the vector $(\log f(X_1; \mu, \sigma^2), \dots, \log f(X_n; \mu, \sigma^2))$. Summing this gives you the log-likelihood function.

Add WeChat powcoder

- ▶ The third method is faster still:

```
nll <- function(params, X){  
  mu <- params[1]  
  sigma2 <- params[2]  
  - sum(  
    dnorm(X, mean=mu, sd=sqrt(sigma2), log=TRUE)  
  )  
}
```

- ▶ `dnorm` stands **d**istribution function of the **normal** distribution (in this case, the value of the density).
- ▶ The optional argument `log=TRUE` returns the log of the density - which is convenient for log-likelihood calculation.
- ▶ This makes your codes succinct and human-readable and idiomatic to the R language.

Assignment Project Exam Help

The negative log-likelihood has been defined. Let us optimise it.

```
optim(par = c(1,1), nLL)
```

- ▶ Above, I just guessed the initial parameter values.
- ▶ Difficult optimisation problems tend to be numerically unstable. That is, if initial parameters are not chosen carefully:
 - ▶ the algorithm might not converge at all, or
 - ▶ might be very slow even when it converges.
- ▶ We want to choose a “good” initial value. One option is to use the Method of Moment Estimator.

```
optim(par = c((mean(X), var(X))), nLL)
```

<https://powcoder.com>
Add WeChat powcoder

Assignment Project Exam Help

- ▶ The previous example seems trivial:
 - ▶ The explicit MLE estimator is known
 - ▶ MLE and Method of Moments estimator coincide.
- ▶ Let's consider a different example. Suppose that X follows a Gamma distribution with parameters (α, β) , i.e.

$$f_X(x) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}, \quad x \geq 0.$$

- We create some dummy data for this task:

```
Y <- rgamma(100, shape=10, rate=10)
```

- ▶ Exercise. Derive the method of moment estimators for α and β . Use them as initial values, retro-fit an MLE estimator on the data Y .

Assignment Project Exam Help

- Say that a variable X follows a Weibull (λ, k) distribution if it has density:

$$f(x; \lambda, k) = \begin{cases} \frac{k}{\lambda} \left(\frac{x}{\lambda}\right)^{k-1} e^{-(x/\lambda)^k}, & x \geq 0 \\ 0, & x < 0 \end{cases}$$

<https://powcoder.com>

- We observe a sample (X_1, \dots, X_n) which are i.i.d. Weibull (λ, k) . Let us estimate λ and k .
- Suppose $n = 100$. Let's simulate some dummy data for this task.

```
dummydata <- rweibull(100, shape=10, scale=20)
```

Add WeChat powcoder

Maximum Likelihood Estimation | Example: Weibull Distribution (Not Examinable) (2)

Assignment Project Exam Help

- Maximising log-likelihood means minimising its negative. The negative log-likelihood can be defined as:

```
negLL <- function(param){  
  shape <- params[1]; scale <- prams[2]  
  -sum(  
    log(scale/shape*(dummydata/scale)^(scale-1)*exp(-(dummydata/scale)^scale))  
  )  
}
```

- R has a shorthand for this. Use the dweibull function to calculate the distribution function of the weibull distribution.

```
negLL <- function(params){  
  -sum(dweibull(dummydata, shape=params[1], scale=params[2], log=TRUE))  
}
```

Assignment Project Exam Help

- ▶ Let's optimise this:

```
optim(par = c(10,10), fn = negLL)
```

- ▶ The output is $\lambda = 10.25363$ and $k = 9.99069$. This is pretty close to the original.
- ▶ How did we choose the initial values?
 - ▶ From other estimation methods e.g. Method of Moments.
 - ▶ Taking "shots in the dark", i.e. guessing.

Add WeChat powcoder